

A Few Notes on Linear Solvers in IC-Ferst

James Percival

AMCG,
Department of Earth Science and Engineering
Imperial College London

IC-Ferst Training 2015

Outline

Direct Methods

Iterative Methods

- Stationary methods

- Krylov subspace methods

Preconditioning

Multigrid

Halting

Parallel Solvers

Solver Failures & Troubleshooting

Numerical Solutions to PDEs

The majority of the work in generating an **algorithm** and set of **discretizations** to solve a system of PDEs numerically is aimed at reducing them to a series of linear equations. I.e. to a **matrix** problem like

$$\begin{pmatrix} 3 & 2 & 0 \\ 4 & 4 & 1 \\ 0 & 2 & 3 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \end{pmatrix} = \begin{pmatrix} 1 \\ 3 \\ 2 \end{pmatrix},$$

or, in a more general form,

$$Ax = b.$$

Solving a linear system

Still actually get a solution for this. IC-Ferst uses linear solvers available from the PETSc framework (<http://www.mcs.anl.gov/petsc/>) to solve problem efficiently and in parallel.

Generally solution algorithms can be split into two sorts, direct methods and iterative methods.

Solving a linear system

1. Direct methods: Generate an explicit representation of A^{-1} and then perform multiplication

$$\mathbf{x} = A^{-1}\mathbf{b}.$$

These methods tend to be slow and memory inefficient, especially when only using the matrix once. Also difficult to do in parallel.

2. Iterative methods: given a previous guess, \mathbf{x}_i , generate a new guess, \mathbf{x}_{i+1} , such that the new residual

$$\mathbf{r}_{i+1} := \mathbf{b} - A\mathbf{x}_{i+1}$$

is “smaller” than the old residual \mathbf{r}_i . Iterate over the algorithm until the residual error is small enough to be acceptable.

A Direct method: LU factorization

As an example of a direct method, let's look at LU factorization. This takes the original matrix, A , and creates two new ones, L & U such that

$$A = LU,$$

where L is lower triangular and U is upper triangular

$$A = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix},$$

$$L = \begin{bmatrix} a & 0 & 0 \\ d & e & 0 \\ g & h - \frac{gb}{a} & \dots \end{bmatrix}, \quad U = \begin{bmatrix} 1 & b/a & c/a \\ 0 & 1 & \frac{ea(f-cd/a)}{ea-db} \\ 0 & 0 & 1 \end{bmatrix}$$

A Direct method: LU factorization

Now $A^{-1} = U^{-1}L^{-1}$ and the triangular matrices can be quickly inverted by two elimination sweeps working from top-to-bottom or bottom to top.

$$L_{11}y_1 = b_1,$$

$$x_N = y_N$$

$$L_{22}y_2 = b_2 - L_{21}y_1,$$

$$x_{N-1} = y_{N-1} - U_{(N-1)N}x_N$$

$$y_n = \frac{1}{L_{nn}} \left(b_n - \sum_{i=1}^{n-1} L_{ni}y_i \right)$$

$$x_n = y_n - \sum_{i=n+1}^N U_{ni}x_i$$

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix}$$

Iterative Methods

Iterative methods can be further subdivided into two groups:

1. Stationary methods: Split A into a bit which is easy to invert and a bit which is moved to the right hand side of the equation.
2. Krylov subspace methods: Get a “good” solution in a subspace smaller than the length of the vector x , then extend the subspace until it spans all of x .

Stationary method: Successive Over-Relaxation(sor)

For this method we use

$$A = D + L + U$$

$$\begin{bmatrix} \text{orange} \end{bmatrix} = \begin{bmatrix} \text{cyan} & 0 & 0 \\ 0 & \text{cyan} & 0 \\ 0 & 0 & \text{cyan} \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 \\ \text{green} & 0 & 0 \\ \text{green} & \text{green} & 0 \end{bmatrix} + \begin{bmatrix} 0 & \text{blue} & \\ 0 & 0 & \text{blue} \\ 0 & 0 & 0 \end{bmatrix}$$

and rewrite the original equation $A\mathbf{x} = \mathbf{b}$ into the form

$$(D + \omega L)\mathbf{x} = \omega\mathbf{b} - (\omega U + (\omega - 1)D)\mathbf{x}.$$

Where ω is a constant, $\omega > 1$. The iterative method is thus

$$\mathbf{x}_{k+1} = (D + \omega L)^{-1} [\omega\mathbf{b} - (\omega U + (\omega - 1)D)\mathbf{x}_k].$$

Knowing the best value of ω to use can be tricky.

Krylov subspace methods: Conjugate Gradients (cg)

If A is a symmetric ($A^T = A$), positive definite matrix ($x^T A x > 0$ for all $x \neq 0$), then the solution to $Ax = b$ is also the value of x which minimises the quadratic form $I(x) = \frac{1}{2}x^T A x - x^T b$.

Could just slide down gradient

$$r_k = Ax_k - b$$

$$x_{k+1} = x_k - \alpha_k r_k$$

Minimizing I implies

$$r_k^T r_{k+1} = r_k^T (Ax_{k+1} - b) = r_k^T (Ax_k - \alpha_k r_k) = 0$$

$$\alpha = \frac{r_k^T A x_k}{r_k^T r_k}.$$

Krylov subspace methods: Conjugate Gradients (cg)

Can do better! Use the new residual to make a search direction “A conjugate” to all the previous ones

$$\mathbf{r}_k = \mathbf{A}\mathbf{x}_k - \mathbf{b}$$

$$\mathbf{p}_k = \mathbf{r}_k - \sum_{i < k} \frac{\mathbf{p}_i^T \mathbf{r}_k}{\mathbf{p}_i^T \mathbf{A} \mathbf{p}_i} \mathbf{p}_i$$

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \mathbf{p}_k$$

with

$$\alpha = \frac{\mathbf{r}_k^T \mathbf{r}_k}{\mathbf{p}_k^T \mathbf{A} \mathbf{p}_k},$$

which makes $\mathbf{r}_k^T \mathbf{r}_k = 0$.




illustration by Oleg Alexandrov -
via Wikimedia Commons

Ksp: Generalized minimum residual (gmres)

Conjugate gradients method relies on A being symmetric and positive (or negative) definite. In the general case where A is non-symmetric but an inverse exists then a similar approach can be used.

1. Work in space $\mathcal{K}_n = \text{span} \{ \mathbf{r}_0, A\mathbf{r}_0, A^2\mathbf{r}_0, \dots, A^n\mathbf{r}_0 \}$,
2. Find $\mathbf{x}_n \in \mathcal{K}_n$ which minimises $\|\mathbf{r}_n\|_{K_n}$.
3. If residual not small enough form
 $\mathcal{K}_{n+1} = \text{span} \{ \mathbf{r}_0, A\mathbf{r}_0, A^2\mathbf{r}_0, \dots, A^{n+1}\mathbf{r}_0 \}.$

In a practical implementation, often impossible to work with \mathcal{K}_n as n gets big. Answer: stop solve periodically and restart with

$$\mathbf{x}_0 = \mathbf{x}_{n_{\max}},$$

$$\mathcal{K}_0 = \text{span} \{ \mathbf{r}_{n_{\max}} \}.$$

Convergence can be slow. Preconditioning is important.

Preconditioning: The Matrix Condition number

A square matrix can be written as

$$A = PDP^{-1}$$

where D is a diagonal matrix containing the eigenvalues of A and P has columns containing the eigenvectors of A . We define the matrix condition number as

$$\kappa := \frac{\max |\lambda_i|}{\min |\lambda_i|} \geq 1.$$

Generally low condition numbers mean a matrix is “easier” to solve numerically, that iterative methods converge and that solutions are more robust to truncation error.

Preconditioning: The Preconditioning matrix

Need to solve $Ax = b$. Suppose we have another matrix M , “close” to A , but easier to invert. Can attempt to solve

$$M^{-1}Ax = M^{-1}b, \text{ (left preconditioning)}$$

$$\text{or} \quad AM^{-1}Mx = M^{-1}b \text{ (right preconditioning)}$$

i.e. solve a two part system

$$My = b,$$

$$AM^{-1}y = b,$$

$$M^{-1}Ax = y.$$

$$Mx = y$$

Advantage is that matrix product is hopefully more “identity-like”, i.e. smaller condition number, so both problems are easier to solve.

Preconditioning: The Preconditioning matrix

Jacobi method

Precondition using $M = D$

SOR


Precondition using $M = D + \omega L$

Incomplete LU factorization

Like a lazy LU factorization. Do trickery to find an approximate \tilde{L} and \tilde{U} such that $\tilde{L}\tilde{U} \approx A$ then use $M = \tilde{L}\tilde{U}$ as preconditioner.


Multigrid

Preconditioning shows it is better to solve simple problems than more difficult ones. Multigrid takes this a step further: Big matrices are hard to solve for, so why not solve a smaller problem which is “like” the big one, and then use the update to correct the full problem?



\Rightarrow (smoothing)

\Leftarrow (prolongation)



Halting Criteria

Good:

Absolute tolerance achieved, $\|r_k\| \leq \tau_{abs}$,

Relative tolerance achieved, $\|r_k\| \leq \tau_{rel} \|r_0\|$,

Bad:

Maximum iterations reached, $k = k_{max}$

Matrix solve diverges $\|r_k\| \geq \tau_{div} \|r_{k-1}\|$,

NANs start appearing

Parallel Solvers

In principle a parallel matrix solve could work just like a serial one. In practice, multiplying out rows of a matrix is local and cheap, while exchanging column information requires expensive communication. PETSc modifies methods to be more efficient. This means

- ▶ Direct solvers don't work
- ▶ Parallel solvers may need more (fast) iterations for good answer.
- ▶ Serial & parallel solves can halt on different answers.

Solver failures

When a PETSc linear solve fails an error message usually gets reported

```
WARNING: Failed to converge.
```

```
PETSc did not converge for matrix solve of: DeltaP
```

```
Reason for non-convergence: KSP_DIVERGED_ITS
```

```
Number of iterations: 3000
```

```
Sending signal to dump and finish
```

This names the variable for which the solve failed, the halting criterion and the number of solve iterations successfully performed.

Solver failures

Possible reasons include

KSP_DIVERGED_ITS Specified limit on solver iterations reached

KSP_DIVERGED_DTOL The residual has increased too much
between iterations

KSP_DIVERGED_NAN The problem PETSc has been asked to solve
has stopped making sense.

KSP_DIVERGED_INDEFINITE_PC You're trying to solve a badly
assymetric matrix using CG.

If you see a much longer error message involving PETSc then either
you're trying to do something it can't (e.g. a direct solve in parallel),
or something else has broken earlier.

Troubleshooting

- ▶ When trying to understand a solver failure, it's important to check that you're looking at the **first** thing which went wrong.
- ▶ If solver failures appear (almost) immediately, then it's possible there's a mistake in your input file.
- ▶ Solver failures can be the result of a bad mesh, either fixed or adaptive. Fix the mesh before trying to fix the solver.
- ▶ If you're **really** sure it's just the solver, then you can try
 - ▶ Increasing the maximum number of iterations
 - ▶ Adding an absolute tolerance limit (e.g. $1.0\text{e-}10$)
 - ▶ Reducing the simulation timestep.

Final Summary

Robust option choice:

- ▶ `Iterative_method(gmres)`
 - ▶ `restarts: 30`
- ▶ `preconditioner(hypre)`
- ▶ `relative_error: 1.0e-7`
- ▶ `max_iterations: 1000`

References



Y. Saad

Iterative Methods for Sparse Linear Systems.
SIAM 2003



G. H. Golub & C. F. Van Loan

Matrix Computations
Johns Hopkins University Press 1996



Fluidity Manual

Applied Modelling & Computation Group



Petsc Manual/online documentation

<http://www.mcs.anl.gov/petsc/>